Next, we will provide several rewriting examples.

To turn on the rewrite engine, use:

RewriteEngine on

To force a redirect to a new location for a document, use:

RewriteRule ^/old.html\$ new.html [R]

To rewrite a **URI** based on the web browser (without a redirect), use:

RewriteCond %HTTP\_USER\_AGENT .\*Chrome.\*

RewriteRule ^/foo.html\$ /chrome.html [L]

To create a human-readable **URI** from a **CGI** script (Note the **PT** or **Passthrough** flag: without it, the rewrite is treated as a file, not a **URI**), use:

RewriteRule ^/script/(.\*) /cgi-bin/script.cgi?\$1 [L,PT]

8.7. Mod Alias

The mod\_alias command maps URIs to filesystem paths outside of normal DocumentRoot. It also provides Alias and Redirect directives, Apache access control considerations, file permission considerations, and SELinux considerations.

Mod\_alias is much simpler than mod\_rewrite. A redirect directive sends a redirect that is similar to the R flag for mod rewrite.

To serve the files located at /newdocroot/ under the URI /new/ (Note the <Directory> stanza for access control modifications), use:

```
Alias /new/ /newdocroot/
<Directory /newdocroot/>
    Options Indexes FollowSymLinks
    AllowOverride None
    Require all granted
</Directory>
```

Care must be taken to properly set the **SELinux**, **Apache** access controls, and filesystem permissions to allow **Apache** to serve content from the directory if it is not inside the default document root.

AliasMatch is like the Alias directive with regular expressions. It is more powerful than Alias, but not as powerful as a rewrite.

To serve the files located at /newdocroot/<SUBDIR> under the URI /new/<SUBDIR> with different permissions for each use:

```
AliasMatch ^{}/new/(.*) /newdocroot/$1/
<Directory /newdocroot/foo/>
    Options Indexes FollowSymLinks
    AllowOverride None
    Require all granted
</Directory>
<Directory /newdocroot/bar/>
    Options -Indexes FollowSymLinks
    Require all granted
</Directory>
...
```



## 8.9. AliasMatch Considerations

AliasMatch is not a direct replacement for Alias. Because AliasMatch uses a regular expression, Alias /new//newdocroot/ is not the same as AliasMatch /new/ /newdocroot/.

The latter will send all **URIs** which contain the string /new/ anywhere to /newdocroot/.

The proper way to replace Alias with AliasMatch is:

```
AliasMatch ^/new/(.*)$ /newdocroot/$1
```

CGI scripts should not be in your document root. There are two ways to run scripts outside your document root:

- Create an alias and set the handler to cgi-script.
- Use ScriptAlias.

Use the same considerations as Alias and AliasMatch:

- · Apache access control considerations.
- File permission considerations.
- SELinux considerations.

Because of the nature of CGI scripts, they should not be directly served as plain files.

ScriptAlias will set up an alias and enable the cgi-script handler. An example of this is:

```
ScriptAlias /newscripts/ /newscripts-docroot/
<Directory /newscripts-docroot/>
Options Indexes FollowSymLinks
AllowOverride None
```

Mod\_Status has the following characteristics:

- · Apache server status.
- · Human-readable page.
- · Machine-readable page.
- · Security considerations with access control and extended status.

To enable mod status, use:

```
<Location /server-status>
    SetHandler server-status
    Require ip 10.100.0.0/24
</Location>
```

To get the machine-readable version of the status page, append auto to the **URI** you have defined above.

Because of the nature of this information, it should be restricted to trusted hosts or networks.

**Apache** version 2.3.6 changed the default for the **ExtendedStatus** directive to be on.

8.12. Mod\_Include

Mod\_Include provides filtering before a response is sent to the client. It is enabled with the +Includes directive. By default, this is enabled only for \*.shtml files, but it can be enabled for all \*.html files with the XBitHack directive.

To enable includes use:

<Location /dynamic/>

```
Options +Includes

XBitHack on

</Location>

An HTML file which had includes would contain
```

An **HTML** file which had includes would contain:

```
<!--#set var="dude" value="very yes" -->
...
<!--#echo var="dude" -->
...
<!--#include virtual="http://localhost/includes/foo.html" -->
```

When you enable the XBitHack directive, **Apache** only looks for includes where the execute permission is set on the file.

Mod\_Perl is a Perl interpreter embedded in Apache. It can provide full access to the entire Apache API, sharing information between Apache processes. CGI scripts can run unmodified.

Some of the downsides to using Mod Perl are:

- Some CPAN modules are not thread-safe, therefore they do not work with the worker MPM (Multi-Processor Module).
- Some Perl methods like die() and exit() can cause performance issues.

To enable mod perl use:

```
Alias /perl /var/www/perl

<Directory /var/www/perl>

SetHandler perl-script

PerlResponseHandler ModPerl::Registry

PerlOptions +ParseHeaders

Options +ExecCGI

</Directory>
```

```
Alias /perl /var/www/perl
<Directory /var/www/perl>
     SetHandler perl-script
     PerlResponseHandler ModPerl::Registry
     PerlOptions +ParseHeaders
     Options +ExecCGI
</Directory>
For a sample Perl script using mod_perl use:
print "Content-type: text/plain \n";
print "mod perl 2.0 rocks! \n";
```

To allow for different performance characteristics, different multiprocessing modules can be used.

The **Prefork MPM** is the default **MPM** on most **Linux** systems. **Prefork** is non-threaded, and is thus safe for libraries and scripts which are not thread-safe. Because each request is handled by a separate process, the **Prefork MPM** uses more memory and resources to handle the same number of connections as the **Worker MPM**.

The **Worker MPM** is a hybrid multi-process, multi-threaded model. Each sub-process spawns threads which handle the incoming connections. Because threads are lighter on system resources than processes, this module has better performance with less resource utilization than the **Prefork MPM**.

The **Event MPM** is based on the **Worker MPM** but it off-loads some of the request handling to shared processes, thus increasing performance yet again.

## The **Prefork MPM** is controlled by the following directives:

- StartServers: Number of child server processes to create on startup.
- MinSpareServers: Minimum number of child server processes to keep on hand to serve incoming requests.
- MaxSpareServers: Maximum number of child server processes to keep on hand to serve incoming requests. Any spare processes above this number are killed.
- MaxRequestsPerChild: Maximum number of requests a child server process serves.
- MaxClients: Maximum number of simultaneous connections to serve. To increase this number, you may also have to increase the ServerLimit directive.

8.2. Learning Objectives

By the end of this session, you should be able to:

- Explore some of the more advanced configurations of Apache, using:
  - Rewrite module.
  - Status module.
  - Alias module.
  - Include module.







**URL** rewriting is a powerful way to modify the request that an **Apache** process receives. Modifications can include moved documents, forcing **SSL**, forcing login, human-readable **URIs**, and redirecting based on the browser.

## 8.4. Rewriting URLs

Because of the power behind mod\_rewrite, rewriting URLs can be complex. The **Apache** documentation is a big help: http://httpd.apache.org/docs/current/rewrite/.

Rewrite rules can exist in the .htaccess files, the main configuration file, or <Directory > stanzas. The .htaccess files seem the most common, but the main configuration file may be more secure.

The mod\_rewrite uses PCRE (Perl Compatible Regular Expressions). This provides nearly unlimited conditions or filters to include, like:

- · Time of day.
- · Environment variables.
- URLs or query strings.
- HTTP headers.
- External scripts.

Use a rewrite map for advanced or dynamic rewrite rules. The **Apache** documentation covers each of them in detail: <a href="http://httpd.apache.org/docs/current/rewrite/rewritemap.html">http://httpd.apache.org/docs/current/rewrite/rewritemap.html</a>.

Pre-compiled rewrite maps exist for specific purposes. To learn more, see the **WURFL** (wireless universal resource file) map: <a href="http://wurfl.sourceforge.net/">http://wurfl.sourceforge.net/</a>.

Rewrite maps are used when rewrite rules are too complex or numerous. There are many different types of maps:

- txt: Flat text.
- dbm: DBM hash (indexed).
- rnd: Randomized plain text.
- prg: External program.



8.16. Configuring Worker

The Worker MPM is controlled by the following directives:

- StartServers: Number of child server processes to create on startup.
- MaxClients: Maximum number of simultaneous connections to serve.
- MinSpareThreads: Minimum number of worker threads to keep on hand to serve incoming requests.
- MaxSpareThreads: Maximum number of worker threads to keep on hand to serve incoming requests.
- ThreadsPerChild: Constant number of worker threads in each child server process.
- MaxRequestsPerChild: Maximum number of requests a server process serves.

8.17. Load Testing

Apachebench (ab) is the tool provided as a part of the Apache httpd server package.

The httperf tool was written at HP Labs. To learn more about it, see <a href="http://www.hpl.hp.com/research/linux/httperf">httperf</a>. Httperf can use log files as a source of URIs to test.

Information about siege, a third-party benchmark tool, can be found at <a href="http://www.joedog.org/siege-home/">http://www.joedog.org/siege-home/</a>.

A tool named **sproxy** works as an **HTTP** proxy server, collecting all the information and **URIs** to use as the **Siege** testing corpus.

The best testing results come from using **URIs** which match the real-life workflow. Some of these tools can use log files directly, others collect information in other ways. Having a proper testing corpus is almost as important as the settings you use.

**Apache** can act as a forward or reverse proxy that is managed by mod\_proxy and mod\_proxy\_http.

Debugging headers include:

X-Forwarded-For

X-Forwarded-Host

X-Forwarded-Server

To enable a reverse proxy to an internal server, use:

```
<Location /foo>
```

ProxyPass http://appserver1.example.com/foo

</Location>

You can also use the following alternative syntax:

ProxyPass /foo http://appserverl.example.com/foo

If the first argument to **ProxyPass** has a trailing /, then the second should have as well.

Apache can also cache requests it proxies, using mod\_cache.

Specialty appliances.

- There are special-purpose tools which do caching much more efficiently:
  - Varnish: a RAM and disk-based cache which uses the Linux kernel's swapping mechanism to speed up caching.

  - Squid: A general-purpose caching proxy.

## 8.19. Speciality HTTP Servers

Specialty **HTTP** servers have been developed to deal with different issues.

The C10k Problem describes the issues in making a web server which can scale to ten thousand concurrent connections/requests.

Many new, lightweight web server systems have been written with the C10k problem in mind:

- cherokee: <a href="http://www.cherokee-project.com/">http://www.cherokee-project.com/</a>
   cherokee has an innovative web-based configuration panel. Benchmarks have shown cherokee to be much faster than Apache for dynamic and static content.
- nginx: <a href="http://wiki.nginx.org/Main">http://wiki.nginx.org/Main</a>
   The nginx server has a small footprint, so it scales well from small servers to high performance web servers. nginx powers some of the largest sites on the Internet, and now accounts for 14% of the web server market share. To learn more, see: <a href="http://news.netcraft.com/archives/2016/02/22/february-2016-web-server-survey.html">http://news.netcraft.com/archives/2016/02/22/february-2016-web-server-survey.html</a>
- Lighttpd
   Lighttpd is used to power some high-profile sites, such as YouTube and Wikipedia. Lighttpd has a light-weight and scalable architecture as well.