13.16. Other Network Filesystems

\$

Other network filesystems include:

- AFS (Andrew File System):
 A distributed, network filesystem built for high performance.
- Gluster:
 A distributed NAS filesystem now maintained by Red Hat.
- GFS (Global File System):
 A shared-disk or clustering filesystem. Multiple nodes can share the same disks.
- DRBD (Distributed Replicated Block Device):
 A distributed shared-disk filesystem build for HA clusters.

The most common errors found in setting up network filesystems are:

- Incorrect firewall settings:
 Older versions of NFS are more difficult to check; you need to open multiple ports to do so.
- Incorrect access control settings:
 Test from a different guest in the same network, or in a different network.
- Syntax errors in configuration files:
 Use testparm, showmount and the like for debugging.
 - NFS: showmount -e <SERVER>
 - SMB: smbclient

13.18. Mounting Network Filesystems

7

The three methods of mounting network file systems (immediate mount - command line-, always mounted, and mounted on demand) use a common configuration file: /etc/fstab. Over the years, additional features and services have updated the options in /etc/fstab. Functionality that required external packages, services and configuration are now combined under systemd services.

The /etc/fstab has slightly changed its purpose with **systemd** from configuration file to drive the **mount** command during system initialization to an input configuration file for the **systemd-fstab-generator** that creates native **unit** files used by **systemd.mount** during system startup. The regular **mount** command still looks in the /etc/fstab file.

Since the preferred method to configure file system mounts is still /etc/fstab, nothing has changed for most mount options. See man systemd.mount for additional details.

13.19. Mount Command

The universal mounting command can mount many types of filesystems, including **NFS** and **cifs** filesystems.

- If the "device" name is in the format SERVERNAME: /share, then NFS is assumed.
- If the "device" name is in the format SERVERNAME//share, then cifs is assumed.
- In most cases, the -t fstype is optional
- Supports options specific to the filesystem, like the cifs username and do-main options.





13.20. Persistent Mounting Network File Systems

Network filesystems can be persistently mounted through /etc/fstab:

- the fs vfstype field must be set, usually to one of the following: nfs, cifs or nfs4
- the fs mntopts field can contain options specific the network filesystem
- the fs_freq field is generally set to 0
- the fs_passno is generally set to 0; setting this field to a non-zero value may cause a delay during system startup time waiting for the mount to complete.

Mounting remote filesystems in this way has the following attributes:

Advantages

- remote filesystem mounted during system initialization
- no connection setup delay when accessing the remote filesystem.

Disadvantages

- constant connection traffic from the server to the client
- network conditions may disrupt the KeepAlive, leaving either the server or the client in an indeterminate state server restarts may require manual client reconnecting.



13.21. Automount Network File Systems

The automount will monitor a configured mount point and, if the information in the filesystem is accessed, the mount will be executed.

There may be an option to disconnect an idle connection. There are a couple of options for automounting:

autofs

- requires additional packages to be installed
- multi-part configuration files
- may require scripts as helpers for mounting some filesystem

systemd.automount

- integrated into systemd
- only available on systemd distros
- unified configuration files, uses /etc/fstab
- uses common systemd.mount facilities, no scripts, but allows unit file overrides.

With the **systemd.automount** implementation, the actual **unit** configuration files are generated from the entries in /etc/fstab. There are options to add to the /etc/fstab to indicate to the **systemd-fstab-generator** the entries are to be automounted. The **systemd unit** files that are generated are located in the /run/systemd/generator directory. The file

names end in **mount** and **automount**. Like other **systemd** files, they can be overridden by entries in the /etc/systemd/system directory, but it is easier to just change the options in /etc/fstab. The **man** page for **systemd.mount** has information on the options. Some examples from /etc/fstab are provided below:

127.0.0.1:/home/export/nfs /home/share/nfs nfs x-systemd.automount,x-systemd.idle-

automounted. The systemd unit files that are generated are located in the /run/systemd/generator directory. The file

timeout=10, noauto, _netdev 0 0 //localhost/cifs-share /home/share/cifs cifs creds=/root/smbfile,x-systemd.automount,x-systemd.idle-timeout=10,noauto,_netdev 0 0

13.2. Learning Objectives

By the end of this session, you should be able to:

- Configure NFS servers.
- · Configure a Samba server.
- Discuss other options for file servers: AFS, Gluster, GFS, and DRBD.



The **N**etwork **F**ile **S**ystem (**NFS**) is a filesystem protocol created by **Sun Microsystems**.

NFS is built upon the Open Network Computing Remote Procedure Call system. RPC connections are managed in Linux using the portmap service.

NFS relies on daemon processes (portmap, nfsd, mountd) and depends on the NFS version used.

In the early versions of **NFS**, security was dependent only on the **UID** or **GID** numbers provided by the client. Later versions allow for different security methods.

13.4. NFS History



NFS version 1 was internal to Sun and was never released publicly.

NFS version 2 was the first publicly available version of **NFS**. It operated over **UDP**, and had no support for files over 2GB in size. Some vendor implementations of **NFS** v2 allowed for **TCP** access.

NFS version 3 introduced the 64-bit file handles (over 2G files), TCP support, and other performance improvements.

NFS version 4 enforces security, has a stateful protocol, and was the first version created after Sun handed development off to the Internet Engineering Task Force (IETF).

The main configuration of the **NFS** server is done in the /etc/exports file. You can reload the server or use exportfs -ra.

Common options include root squashing, read size, write size, and read/write.

The syntax of the exports file is:

<DIRECTORY> <HOST OR NETWORK>(<OPTIONS>)

Sharing the /srv/nfs/ directory would have an entry like the one below:

/srv/nfs/ 192.168.122.0/24(rw,sync,root_squash)

The NFS client "mounts" the remote filesystem on to the local system. There are a couple of commands specific to NFS:

• **showmount** command: Queries the mount daemon on the remote server for information, including the "shares" that are available for mounting:

```
# showmount -e SERVER
Export list for SERVER:
/srv/exports *
```

- portmap running: A dynamic port mapping daemon init ally designed to reduce the usage of well-known port numbers; still used by NFS up to version 3. Not required by NFS version 4; it is required for the showmount command.
- mount command: The mount command has the filesystem type NFS, which links to the mount.nfs command. There
 are two formats of the mount command for NFS shares.
 - One option for mounting NFS shares is:

mount HOST:/export /mount-point

where the host:/export portion causes the mount command to process this mount as NFS.

- The other form of the mount command is:

- portmap running: A dynamic port mapping daemon init ally designed to reduce the usage of well-known port numbers; still used by NFS up to version 3. Not required by NFS version 4; it is required for the showmount command.
- mount command: The mount command has the filesystem type NFS, which links to the mount.nfs command. There
 are two formats of the mount command for NFS shares.
 - One option for mounting NFS shares is:
 - mount HOST:/export /mount-point

where the host:/export portion causes the mount command to process this mount as NFS.

- The other form of the mount command is:
 mount -t NFS HOST:/export /mountpoint
- which specifies NFS is being used.
- # mount SERVER:/share /mnt/share

Note that both examples above use the construct **server**:/**sharename** for mounting. The *sharename* is discovered with the **showmount** -e **servername** command.



13.7. NFS Security Considerations

The NFS default security is to use the UNIX UID and GID. The root_squash option translates the root user's UID/GID (0) to an anonymous UID/GID. You should not disable root squash without a good reason.

There are other authentication options available for **NFS**. The **Kerberos** authentication is commonly used to overcome the security issues of **UID/GID** mappings.

13.8. NFS Performance Considerations

Many factors contribute to the speed and performance of an NFS server or client.

Properly setting the values of rsize and wsize will allow for greater speed in a file transfer. However, you can only reasonably increase the block size to the MTU of your network between client and server. Increasing the frame size (Jumbo Frames) is one option.

Moving from a 1G to a 10G Ethernet network would vastly speed up an NFS setup.

The asynchronous mode trades speed for lack of robustness. An unclean shutdown of a server or client operating in an asynchronous mode has the potential to corrupt the data.



13.9. SMB/CIFS Overview

Ŧ

The Server Message Block (SMB) protocol was originally designed at IBM and later incorporated as the de-facto networking file/print sharing system for Microsoft Windows.

In 1996, the latest version of the **SMB** protocol was renamed **CIFS** (Common Internet File System), as many new features were added.

The **Samba Project** started as a reverse-engineered implementation of the **SMB** protocol for **Solaris** servers. To learn more about this, visit http://www.samba.org/.

Samba is built to be an SMB/CIFS server which will run on any UNIX-like system.

13.10. Samba Features

Samba features include the following:

- · Samba can create file or printer shares.
- Samba version 3.x can act as a WindowsNT domain controller.
- Samba version 4.x can act as an Active Directory domain controller.
- Samba version 4.x is available on most distributions.

The default location for the **Samba** configuration file is /etc/samba/smb.conf, which uses an **INI**-file-like syntax, with section headers enclosed in square brackets.

```
[global]
  workgroup = MYGROUP
  server string = Samba Server Version %v
  log file = /var/log/samba/log.%m
  max log size = 50
  cups options = raw
```

Each individual share then goes into its own section:

```
[mainshare]
    path = /srv/exports/
    read only = yes
    comment = Main exports share
```

Due to the difference in system password hashing mechanisms, **Samba** cannot verify some user passwords without additional help.

The **smbpasswd** command allows you to manage your passwords in both the **Samba** password file and in directory services.

To create a new password entry for Samba, do:

smbpasswd -a geoff

To change the password for **geoff**, do:

smbpasswd geoff

In the event that the **UNIX** username does not match the **Samba** username, the /etc/samba/smbusers file allows for the translation.

UNIXNAME = SMBNAME SMBNAME2



13.13. Testing the smb.conf Syntax

Use the following stanza to test the smb.conf syntax:

```
server string = Samba Server Version %v
log file = /var/log/samba/log.%m
max log size = 50
cups options = raw
[mainshare]
comment = Main exports share
path = /srv/exports/
```

workgroup = MYGROUP

[global]

Once you have created your smb. conf file, test the syntax with the testparm command:

```
# testparm
Load smb configuration files from /etc/samba/smb.conf
rlimit_max: increasing rlimit_max (1024) to minimum Windows limit
(16384)
Processing section "[mainshare]"
```

Server role: ROLE_STANDALONE

Press enter to see a dump of your service definitions

When you press enter, you will get a parsed copy of the configuration file.

Loaded services file OK.

13.14. Samba Clients

There are several ways to interface with a **Samba** server:

Query the shares on a server:

LFLAB

\$ smbclient -L 172.16.104.131 -U student

Enter SAMBA\student's password:

Domain=[UBUNTU] OS=[Windows 6.1] Server=[Samba 4.5.8-Ubuntu]

Sharename	Type	Comment
cifs-share	Disk	Example share for testing home-export-cifs
TPCS	TPC	TPC Service (ubuntu server (Samba, Ubuntu))

Domain=[UBUNTU] OS=[Windows 6.1] Server=[Samba 4.5.8-Ubuntu]

UBUNTU

Server	Comment	UBUNTU
Workgroup	Master	

- Using the FTP-like smbclient command:
- \$ smbclient //SERVER/mainshare
 smb: \> get /foo
- Mounting the **SMB/CIFS** share into your name space:

```
# mount -t cifs -o username,password //SERVER/share/ /mnt/point/
```

To avoid putting usernames and passwords into a world-readable file (/etc/fstab), the CIFS mount command takes the credentials=filename option.

```
# mount -t cifs -o credentials=filename //SERVER/share/ /mnt/point/
```

The credentials file has the following syntax:

username=value password=value domain=value 13.15. Samba Tools

The distribution-specific tool for **CentOS**, **system-config-samba**, has been deprecated and removed. The **Samba** organization has also deprecated and removed the included web administration tool **SWAT**.

Several other graphical interfaces for **Samba** are available. For more information, see http://www.samba.org/samba/GUI.

One popular tool is the **webmin** tool, available from http://www.webmin.com.

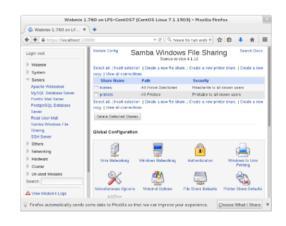


Figure 13.1: Webmin Samba