# By the end of this session, you should be able to:

- Examine topwrappers as an application firewall.
- Discuss the netfilter firewall.
- · Configure the netfilter firewall with iptables.
- Examine available options for managing firewalls.
- Discuss a forced transparent proxy.



15.3. TCP Wrappers

The **TCP Wrappers** system is a host-based network firewall and ACL. Originally, it only protected the **inetd** system, but has now been extended with the shared object library **libwrap**.

The configuration for **tcpwrappers** is handled by two files, /etc/hosts.allow and /etc/hosts.deny. Both files have the same syntax:

### <DAEMON>:<CLIENT>

The <DAEMON> should match the name of the binary of the service (e.g. sshd). The <CLIENT> pattern can be:

- An IP address: 10.30.21.7.
- A network/netmask: 10.30.21.0/255.255.25.0.
- A domain name: .foo.example.com.
- A partial address: 10.30.21.
- A file name full of the above patterns: /etc/ssh-hosts.allow

When traffic comes to a **libwrap**-enabled **daemon**, those two files are consulted to see the following:

When traffic comes to a **libwrap**-enabled **daemon**, those two files are consulted to see the following: If the pattern matches in /etc/hosts.allow, the traffic is permitted.

• If the pattern is not found in /etc/hosts.allow, and it matches in /etc/hosts.deny, the traffic will be denied.

• If the pattern does not match in either file, the traffic will be permitted.

15.4. TCP Wrappers Examples

7

Below you will find some examples of **TCP wrappers**:

hosts.allow:

vsftpd:ALL

ALL:LOCAL

ALL:10

ALL:.example.com EXCEPT untrusted.example.com

hosts.deny:

ALL:ALL

**netfilter** is a packet-filtering framework built into the **Linux** kernel. To better understand **netfilter**, we need to start with some vocabulary:

- . The netfilter firewall consists of Tables.
- Tables consist of Chains.
- Chains have a default Policy.
- Chains consist of Rules.
- Rules consist of a Match criteria and a Target.

Rules in each chain are read first to last, and the first match wins. If a packet does not match any rule in a chain, the policy of the chain applies.

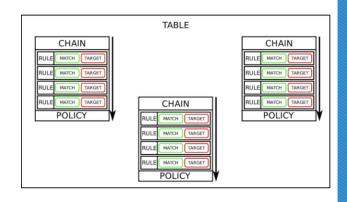


Figure 15.1: netfilter Vocabulary

The **filter** table deals with packets bound for the local machine, being routed through the machine, or packets generated by processes on the machine. It contains the default **chains**:

- INPUT
   For packets bound for local processes.
- FORWARD
   For packets being forwarded through the machine.
- OUTPUT
   For packets generated by local processes that and are now outbound to the network.

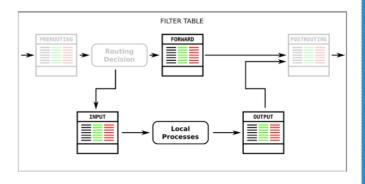


Figure 15.2: Filter Table

## The Network Address Translation

(NAT) table is used when traffic that creates a new network connection is encountered. It contains the default **chains**:

#### PREROUTING

For altering packets just as they come in.

#### • OUTPUT

For altering packets generated by local processes, prior to routing.

#### • POSTROUTING

For packets just about to go out to the network.

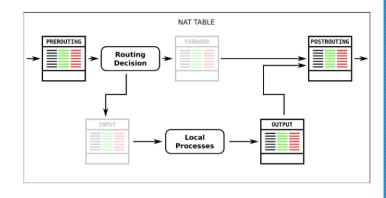


Figure 15.3: NAT Table



15.8. Mangle Table

The **mangle** table is for specialized manipulation of network packets. It contains the following default **chains**:

### • INPUT

For packets bound for local processes.

#### PREROUTING

For altering packets just as they come in.

### FORWARD

For packets being forwarded through the machine.

#### • OUTPUT

For packets generated by local processes that are now outbound to the network.

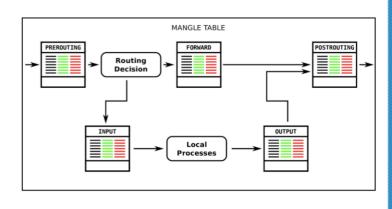


Figure 15.4: Mangle Table

- OUTPUT
   For packets generated by local processes that are now outbound to the network.
- network.

   POSTROUTING
  For packets just about to go out to the

network.

15.8. Mangle Table

The **mangle** table is for specialized manipulation of network packets. It contains the following default **chains**:

### • INPUT

For packets bound for local processes.

### • PREROUTING

For altering packets just as they come in.

### FORWARD

For packets being forwarded through the machine.

### • OUTPUT

For packets generated by local processes that are now outbound to the network.

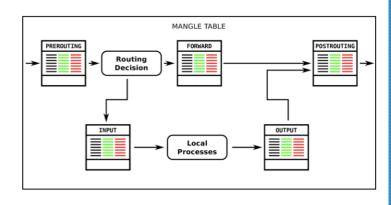


Figure 15.4: Mangle Table

 POSTROUTING For packets just about to go out to the Figure 15.4: Mangle Table network.

## The **iptables** command can be broken into multiple pieces:

- The table using the -t switch. If no table is specified, the default is filter.
- The command, which is one of the following:
  - '-I' (insert) Create a new rule at the top of the chain.
  - '-I #' (insert) Create a new rule at position # of the chain.
  - '-A' (append) Create a new rule at the bottom of the chain.
  - '-P' (policy) Change the chain's policy.
  - '-D' (delete) Delete a rule.
  - '-D #' (delete) Delete rule number #.
- · The chain name.
- The match criteria.
- The target.

15.9.b. Iptables Command (Cont'd)

```
iptables -t filter -A INPUT -m tcp -p tcp --dport 22 -j ACCEPT
```

This rule appends the rule to the bottom of the INPUT chain, loads the top module, matches the TCP protocol destination port 22 and jumps to the ACCEPT target.

Here is another example:

Insert a new rule at the top of a chain:

```
# iptables -I INPUT -m udp -p udp --dport 53 -j ACCEPT
```

Set the INPUT chain policy to 'DROP':

```
# iptables -P INPUT DROP
```

Delete the third rule from the INPUT chain:

```
# iptables -D INPUT 3
```



15.10. Match Criteria

V

The match criteria is the essence of the **iptables** system, and allows for a lot of flexibility. By default, if you use the -p or -protocol switch, the corresponding module is loaded. Some of the more common modules used are the following:

- state:
  - Allows for matching for stateful firewalls.
- tcp:
- Allows for matches on the **TCP** information:
- Source port.Destination port.
- top flags.
- udp:
- Allows for matches on **UDP** information:
- Source port.
- Destination port.
- icmp:

Allows for matches on ICMP query types.

• icmp:
Allows for matches on ICMP query types.

A full listing can be found in the iptables man page, under the heading MATCH EXTENSIONS.

15.11. Targets

•

The default targets in the **netfilter** system are:

• ACCEPT

Pass the packet along to the next stage.

• DROP

Send no response to this packet and ignore it.

• RETURN

Go back to the calling CHAIN and start processing on the next rule.

REJECT

Send a message back explaining why the packet is not allowed.

You can also define your own chain, which you can then use as a custom target.

15.12. Configuration Utilities

₹\_

Because of the inherent complexity of the **netfilter** management, many tools have been created to help alleviate the burden on systems administrators. The **CentOS system-config-firewall**, the **Ubuntu gufw**, and the **OpenSUSE yast firewall** are examples of configuration utilities.

Each distribution has its own GUI or TUI mode firewall tool.

There are also generic tools like shorewall, which wrap the complexity of iptables/netfilter in an API.

The latest addition to firewall management is **firewalld**. The **firewalld** tool is available in most of the recent distributions.

The system-config-firewall utility is a CentOS-specific tool that can be run in either a GUI or TUI mode.

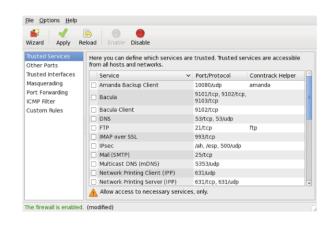


Figure 15.5: System-Config-Firewall

As an **Ubuntu**-specific tool, **gufw** is a **GUI** wrapper of the command-line tool **ufw**.



15.14. gufw

Figure 15.6: gufw



15.15. YAST Firewall



As an **OpenSUSE**-specific tool, the **YAST firewall** runs in a **GUI** or **TUI** mode, depending on how it is started.

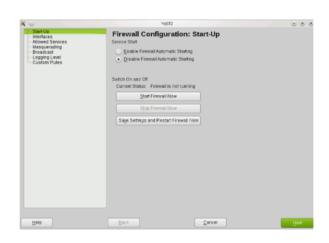


Figure 15.7: YAST Firewall

ŧ

The **firewalld** utility provides a dynamically managed firewall:

- It is available on most distributions.
- It has a GUI and command line interface.
- Configuration changes are applied dynamically.
- Includes support for IPv4, IPv6, and Ethernet bridges.

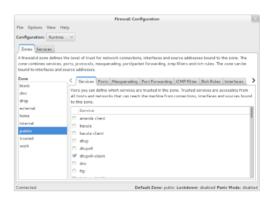


Figure 15.8: firewalld GUI

## 15.17. Distribution Defaults

Each distribution has its own method for storing the iptables state. Some make it easy to manage it with generic tools, while some provide easy-to-use tools for firewall management:

#### On CentOS:

- Firewall rules are saved in /etc/sysconfig/iptables.
- Easiest management is done with **iptables** and **service iptables save**.

### On OpenSUSE:

- Firewall rules are saved in /etc/sysconfig/scripts/SuSEfirewall2-\*.
- Easiest management is done with yast firewall.

### • On Ubuntu:

- The default firewall is managed with ufw.
- Custom rules are stored in /etc/ufw/\*.rules.
- firewalld.

If you manage a server via **SSH**, blocking **SSH** with an **iptables** rule makes management hard. The same goes for setting a restrictive policy without first enabling **SSH** remote access.

- · Set an at job to reset the firewall to allow access while you work:
  - # at now +30 minutes
  - # at> service iptables stop

Another option is to always use a management system to change the firewall rules. At a minimum, use a script to set the rules.

A minimum firewall should allow returning network traffic access, as well as access to the localhost device:

- # iptables -A INPUT -m state --state=ESTABLISHED,RELATED -j ACCEPT
- # iptables -A INPUT -i lo -j ACCEPT

15.20. Forced Transparent Proxy

Using both **iptables** and **Squid**, you can force a network through a proxy server.

The configuration for **Squid** should look like this:

```
httpd_accel_host virtual
httpd_accel_port 80
httpd_accel_with_proxy on
httpd_accel_uses host header on
```

The iptables rules should look like this:

This forced transparent proxy is not entirely secure and is not entirely transparent.

Network Address Translation (NAT) allows for multiple network hosts to share the same external IP address. There are two types of outbound NAT or source NAT:

- MASQUERADE Works with a dynamic source IP address. It is useful for servers with dynamic IP addresses.
- SNAT Works with a static source IP address. It is less complex than MASQUERADE.

There is also a form of inbound or destination **NAT** (**DNAT**). **DNAT** allows for services to be behind a bastion host and to be easily load-balanced to different hosts.

To enable any of these types of **NAT**, the **ip\_forward** kernel option must be set to **1**.

```
# echo 1 > /proc/sys/net/ipv4/ip forward
```

It is also a good idea to make this change in the sysctl.conf file.

An example of a masquerade rule is the following: